

Ch041 Introduction :-

- * purpose of database systems:
 - reduce difficulty in accessing data & programs built for storing data
 - Atomicity of updates {datastructures built}
 - concurrent access by multiple users
 - security problems

+ Schema & Instances

}
an entry in table.

Snapshot of table at some instance.

→ Data definition language DDL

- * for defining schema of database.

```
create table instructor (
    ID      char(5) fixed length!,
    name    varchar(20),
    dept_name varchar(20),
    Salary   numeric(8,2))
```

we have

- Schema
- Integrity constraints
 - primary key
- Authorization
 - who accesses what.

→ Data manipulation language DML (or query language)

- * Basically two types:

Procedural DML: user specifies what data & how to get it.

Declarative DML: user says what data but doesn't specify how...

}
 just declare
 what you want.
 NO involved code.

- * Database schemas: physical, logical.

Database Design: physical, logical

Database Engine:

Storage Manager

Query Processor

Transaction Management Component

To optimize our queries

Ensures database remains in a consistent state; even while power failure

* Architecture:

- Centralized Database
 - Client-server
 - Parallel databases
 - Distributed databases

1921 - Springwood - Chipping Norton -

2 numbers of *Staphylococcus*

Chloride 11

(19) *red* *yellow* *green*

(1987) 10: 49-60, 62-63, 65-66

(03) 200-1133

As the author of the book goes around, he is asked to sign his name.

left out ~~unseen~~

1997-1998 Annual Report of the Joint Select Committee

ch2 : Intro to relational Model:-

relation = Table.

- * Each table consists of columns (Attributes)
& rows (Tuples)

- special value "null" is member of every domain.

instance means,
snapshot of table's
all tuples.

→ Keys :-

- * let $K \subseteq R$. K is Superkey of R if values of K are sufficient to identify unique tuple in R .
 $\{ID\}$ & $\{ID, \text{dept_name}\}$ are both superkeys for Instructor.
- * Superkey K is candidate key, if its minimal.
 - One of the candidate keys is chosen as primary key.
- * foreign key constraint: value in one relation should appear in another relation.
 - : referencing relation
 - : referenced relation

→ Relational Alzebra:

- * This is not turing-machine equivalent. 6 basic operations.
 - A procedural language, consisting of set of operations that take one or two relations as input, & produce new relation as output

select : σ $\sigma_{\text{dept_name} = \text{"Physics}}(\text{instructor})$

project : Π $\Pi_{A_1, A_2, \dots, A_k}(r)$

union : \cup

Set difference: -

cartesian product: \times $r \times s = \sigma_{\text{JOIN}}(r \times s)$

rename: ρ

$\rho_x(E)$

$\rho_{x(A_1, A_2, \dots, A_k)}(E)$

* Assignment operator: temporary relational variable

Physics $\leftarrow \overline{\text{dept_name}} = \text{"Physics"}$ (Instructor)

Ch3 - Introduction to SQL:-

IBM developed SQL

→ Data definition Language:- (of SQL)

- * SQL's DDL allows a specification for

- schema for each relation
- value types for each attribute
- integrity constraints
- set of indices to be maintained for each relation.
- authorization for each relation.
- physical storage structure of each relation.

1. Create tables:-

create table instructor (

```

ID      varchar(5),
name    varchar(20) not null,
dept_name varchar(20),
Salary   numeric(8,2),

```

primary key (ID),

foreign key (dept_name) references department);

end of SQL query.

2) Update to table Schema:-

• Insert into Instructor values ('10211', 'Smith', 'Physics', '66000');

• delete from Students where CPI < 6;

• drop table instructor;

• Alter table instructor add

weight int(3,2);

• Alter table demons drop name;

Basic query

Select a₁, a₂, ...

From r₁, r₂, ...

Where

r₁.name = r₂.name and

* SELECT clause.

SELECT DISTINCT for distinct values.

SELECT ALL a_1, a_2, \dots from ...
to allow duplicates

* can have arithmetic.

> Select runs * 0.1 as avg_runs
over $\Sigma 0.1$

* WHERE clause.

$<$, $>$, \leq , \geq , \neq ; and, or, not
not equal to

> WHERE cond₁ and
cond₂ and
:
cond_K

* RENAME using 'as'; for attributes and relations.

→ String operations

* We have two string matching operators:

(%) matches substrings

(-) matches single character.

where name LIKE '%dar%'

Note:-

i) BETWEEN comparison operator :-

where Salary between 9000 and 10000;

i.e. $9000 \leq \text{Salary} \leq 10000$

Second query: Second oldest result

Third query: Third oldest result

→ Set Operations On Relations:-

() () ()
union intersect except
() () () } automatically
 eliminates duplicates.

For relationship between
multiple relations available with
union, intersect, except
to retain duplicates.

use
union all
intersect all
except all
to retain duplicates.

- * IS NULL is used to check if attribute value is null.

Eg: WHERE Salary is null and
Name is not null;

→ Aggregate Functions:-

- * These operate on multiple values of a column.

> avg, min, max, sum, count (even more exist...
) first, last, rank, dense_rank,

Eg: Select count(*)
From teaches;

Select count(distinct ID)
From instructors;

→ Group by: *Awesome guy!

- * Attributes in SELECT clause; without aggregate function must appear in GROUP BY clause.

Eg: SELECT dept_name, ID, avg(salary)
FROM teaches

GROUP BY dept_name, ID

1) HAVING clause:-

Select deptname, avg(salary) as salary-avg, IMP
from instructor

group by dept-name

having salary-avg > 10,000;

different from 'WHERE' clause

* predicates in WHERE clause are applied

Before forming relation

those in HAVING clause are applied

After forming relation.

2) Nested Subqueries:-

i) in Select clause:-

is replaced by subquery that generates single value.

ii) in From clause:-

we write subquery in ('')s.

iii) in Where clause:-

where salary in (...) and

salary not in (...);

3) Set Membership:-

- where course_id in (select...
from...
where...);

- where course_id not in (
(...));

4) Set Comparison:-

i) SOME clause:

where salary > somes (,);

Sub
query

iii) ALL clause:

where salary > all (---);

Subquery

5) Test for empty relations:-

- * EXISTS clause returns true if argument is not empty.

exists r $\Leftrightarrow r \neq \emptyset$

not exists r $\Leftrightarrow r = \emptyset$

useful for queries like "students who have taken all courses of maths!"

* UNIQUE clause is a test for absence of duplicates.

where UNIQUE (...); is true; If (...) relation has no duplicates.

→ Subqueries in FROM clause :-

* Select dept-name, avg-salary

From (Select dept-name, avg(salary)
from instructor
group by dept-name)

as dept-avg(dept-name, avg-salary)

WHERE avg-salary ≥ 40,000;

* WITH clause : * Good one

- provides a way of defining a temporary relation, visible only to query in which its written

WITH dept-total(dept-name, value) as

(
|
|
) ,

usable immediately.

dept-avg(value) as

{ Select avg(value)
from dept-total }

} also part
of with clause.

SELECT * FROM dept-average;

→ Scalar Subquery:

* used where single value is expected. runtime_error if more than one

Eg: select dept-name,

{ Select count(*) From instructor } Scalar subquery.

where department.dept-name = instructor.dept-name)

as num_instructors } renaming.

from department;

* All we've done is ask, ask, ask. Time to modify the database.

1) Deletion of tuples:-

delete

from instructor

where salary < (select avg(salary) from instructor);

scalar subquery.

computed once per row.

only once.

(common step in update)

2) Insertion of tuples:-

* insert into course

values ('CS-347', 'Comp.Sci', 4);

* insert into course (credits, course_name, dept_name)

values (4, 'CS-347', 'Comp.Sci');

* We can insert a relation into a table.

INSERT INTO instructor

NO VALUES()

SELECT * FROM instructor
WHERE Salary > 10,000

Keywords

3) Update tuples:-

* UPDATE instructors

SET Salary=Salary*1.05

WHERE

salary < (select avg(salary) from instructors);

Q) Increase salaries of people whose salary over 100k by 3%, and all others by 5%.

A) we can write two SQL queries... but CASE clause is more appropriate.

* case statement for conditional update:

(same question as before)

UPDATE instructor

```
SET Salary = CASE
    WHEN Salary <= 100,000 THEN Salary * 1.05
    ELSE Salary * 1.03
END;
```

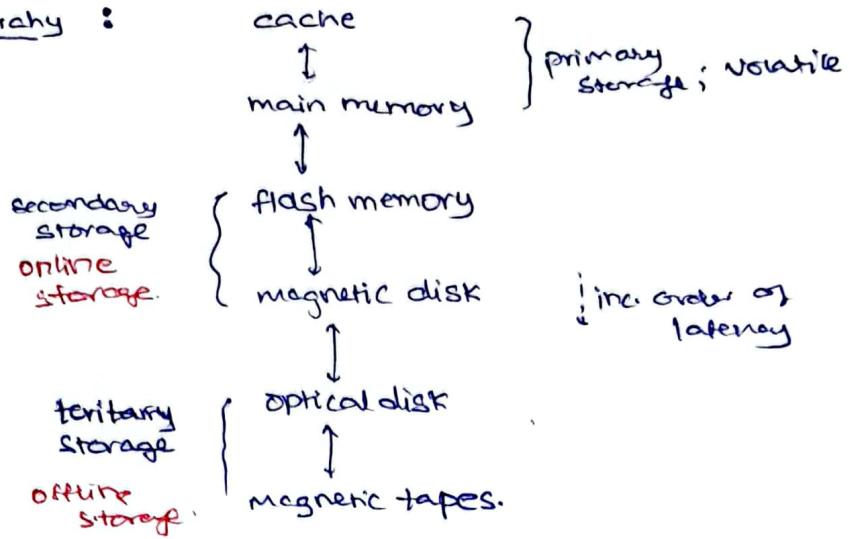
* instead of sum(credits); use

```
case
when sum(credits) is not null then sum(credits)
else 0
end;      to set total-creds to 0, when its null.
```

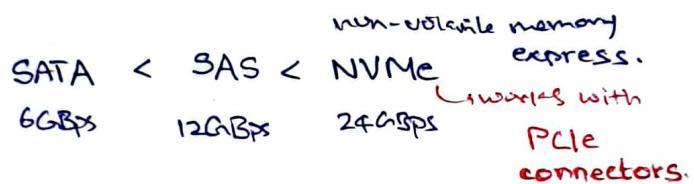
about forms and functions

12. Physical Storage Systems:

* Storage hierarchy :



* Disk interface speeds:

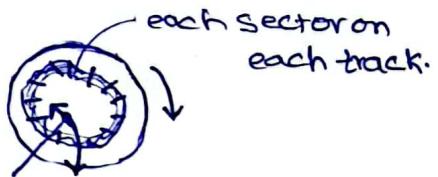


→ Magnetic disks:-

• Access time:

Seek time: reposition arm over correct track.

- Avg. time is $\frac{1}{2}$ th of worst time



Rotational latency:

go to the correct sector.

- Avg. latency is $\frac{1}{2}$ th of worst time.

→ Performance measures:-

1) Disk blocks:- logical unit for storage allocation and retrieval.

- Smaller: more disk operations means

- larger : more space wasted. means

2) Sq. Access pattern:-

- successive req. are for successive disk blocks.

- disk seek required only for 1st block.

3) Random access pattern:-

- successive requests are random.

- Each access needs a seek.

4) IOPS: I/O per second

5) MTTF: avg. time; disk is expected to run continuously.

$$= 1,200,000 \text{ hours} \dots \text{for new disks}$$

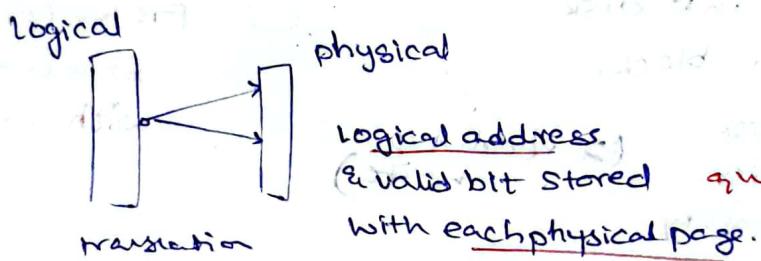
→ Flash storage

* Not much difference b/w sq. and random read.

- pages can be written only once; must be erased to re-write.
- remapping logical page addr. to physical page addr. avoids waiting for erase

Flash translation table tracks mapping

→ Also stored in label field of flash page. (But why!)



→ RAID:

redundant arrays of independent disks

→ disk organization technique; to manage a large number of disks:

- high capacity & speed, using disks in parallel.
- high reliability by storing data redundantly.

ES) mirroring/shadowing: duplicate every disk.

mean time to date loss; depends on MTTF & mean time to repair.

$$\text{MTDL} = \frac{(\text{MTTF})^2}{2 \cdot \text{MTTR}}$$

- management of failure of disks (2 or 3)

* Block Striping:-

n disks; block i goes to $i \text{ mod } n + 1$. not any redundancy.

Raid 10:

- non-redundant
- block-striping



used in high perf.
when safety is not priority.

Raid 11:

- mirrored disks
- block striping.



best write performance.

Parity blocks:

Store XOR of disks.

→ When new entry into disk;
recompute parity block.

- use old disk (2 reads + 2 writes)
- use new disk. (2 reads + 2 writes)
- use old parity.

(here only 1 read + 2 writes)

↓
Prefer for

Small size storage;
with small random updates.

Raid LS: block-interleaved distributed parity.

parity blocks; stored across all N disks; rather than single disk.

so N parallel

in block read/write.

(nice)

D ₀	D ₁	D ₂	D ₃
(P ₀)	0	1	2
3	(P ₁)	4	5
6	7	(P ₂)	8

→ Prefer for large amounts of data storage;

writes sequential.

Raid 2B: stores 2 partitions; to account for multiple disk failures.

* optimization of disk block access:-

Buffering,
prefetching.

Disk arm scheduling to minimize movement.

Cache → Main memory → Flash → Mag disk → opt disk → Tape
 Primary Secondary

Platter has tracks. Tracks has sectors.

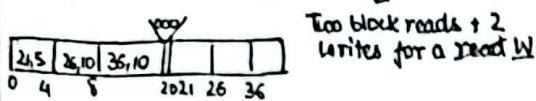
Seek Time: Avg = $\frac{1}{2}$ Worst } Same for Rotational

Wear leveling → done because erase makes storage unreliable

$$MTTDL = (MTTF)^2 / 2MTTR = \frac{1}{n(n+1)} \lambda^2$$

RO: Striping RI: RO + Mirroring

RG: P+Q Redundancy



Heap File: Free-space map → Free_fn = Value / n

Sequential: Ordered by search key

Ch-14 Indices Ordered I: stores search key in sorted order

Clustering: S.K. defines sequential order
Non-clus.: diff. from seq. order

Dense: For every S.K. else sparse

B+Tree: $P_1 K_1 P_2 \dots P_{n-1} K_n P_n$

$\lceil n/2 \rceil$ and $\lceil \rceil$ children for non-leaf

$\lceil (n-1)/2 \rceil$ and $(n-1)$ values for leaf

Complexity: $T \log \lceil n/2 \rceil (N)$

* Bulk Loading → 1. Sort entries first
2. Bottom up B+ Tree constr.



Buffer Tree → Drawback: More random I/O

→ less overheads on query

* K-d trees → partitions 3
one dim into 2.

R-trees → [BB₁, BB₂, BB₃]
[ABIC] [DEF] [GHI]

** Pipelining

open() close() next()
file scan: Initialise output next tuple

Merge join:
soft relations demand driven (lazy)
Producer driven (eager)

Name - Akash G
Roll No - 190050038

Linear Search $t_c + b_r + t_T$

Clustering, Eq. $(hi+1) * (t_T + t_S)$

Clustering, N-K $hi * (t_T + t_S)$
+ $t_S + b * t_T$

Secondary, Eq. $(hi+1) * (t_T + t_S)$

Secondary N-K $(hi+n) * (t_T + t_S)$

n = no. of records fetched

Comparison == Eq. on N-K.

Nested Join → Transfer

$$b_T = n_T * b_S + b_R$$

Seek ← $b_S = n_T + b_T$

Block-Nested $T = b_T * b_S + b_R$

$$S = 2b_T$$

↳ Best-Case → $b_T + b_S : T$

2 seeks

Merge Join $T: b_T + b_S, * b_B = \text{buffer}$

$[b_T/b_B] + [b_S/b_B] \leftarrow \text{seeks}$ blocks for each reln.

Sorting:

$$* BT \rightarrow b_T (2 \times \lceil \log_{M/b_B} \lceil b_T/M \rceil \rceil - 1) + 1$$

$$S \rightarrow 2 \lceil b_T/M \rceil + \lceil b_T/b_B \rceil (2 \times \lceil \log_{M/b_B} \lceil b_T/M \rceil \rceil - 1)$$

* Pre-Midsem

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

rank() over (partition by dept-name order by GPA desc).

not by me

Query Optimisation

$$(E_1 M_{B_1} E_2) M_{B_2} M_{B_3} E_2 \equiv E_1 M_{B_1, B_2, B_3} (E_2 M_{B_2} E_3)$$

$$T_{B_0}(E_1 M_B E_2) \equiv (T_{B_0}(E_1)) M_B E_2$$

$$T_{B_1, B_2}(E_1 M_B E_2) \equiv (T_{B_1}(E_1)) M_B (T_{B_2}(E_2))$$

$$T_{L_1, U_{L_2}}(E_1 M_B E_2) \equiv (T_{L_1}(E_1)) M_B (T_{U_{L_2}}(E_2))$$

$$\equiv T_{L_1, U_{L_2}}(T_{L_1, U_{L_3}}(E_1) M_B T_{L_1, U_{L_4}}(E_2))$$

L_3 = attr of E_1 in θ but not in L_1, U_{L_2}

$2(n-1)!/(n-1)!$: Different join orders

3^n Time : DP Space : 2^n

$n2^n$ Time left-deep Space : 2^n

opt. cost budget : Stop opti. early

Plan caching : reuse prev. plan

$$T_{ASV}(r) \rightarrow n_r \times \frac{v - \min(A_i, r)}{\max(A_i, r) - \min(A_i, r)}$$

Selectivity = S_i/n_r

$$\text{Conj} \rightarrow n_r \times \left(\frac{S_i}{n_r}\right)^n \quad \text{Disj} \rightarrow n_r \left[1 - \left(\frac{S_i}{n_r}\right)^n\right]$$

$$RNS = \emptyset \quad TMS = T \times S \quad G_A(r) = \begin{cases} \text{key of } R & \text{SS} \\ \text{use } n_r * n_s / V(A_i, S) \end{cases}$$

$TMS \rightarrow$ if all attr. are from r
 $V(A_i, TMS) = \min(V(A_i, r), n_r n_s)$

if A contains A_1 from r , A_2 from s
 $= \min(V(A_1, r) * V(A_2 - A_1, s))$, unique values
 $V(A_1 - A_2, r) * V(A_2, s), n_r n_s$

$$\min/\max = \min(V(A_i, r), V(A_i, S))$$

* Transactions / additional Opt.

Semijoin $\rightarrow r \wedge s$, Anti-SJ : $r \overline{\wedge} s$

Select A from r_1, \dots, r_n where $P_1 \in$
 $\exists s_1 \dots s_m$ such that $P_2 \in P_1$

$$= T_R(P_1, r_1, \dots, r_n) K_{P_2} T_{P_1}(s_1, \dots, s_m)$$

involves correlation variables

Material views: (Join) $V^N = V^0 \cup (r \wedge s)$

$$V^N = V^0 - (d \wedge D)$$

For projection,

Keep count changes persist even if system fails

Atomicity / Consistency / Isolation / Durability

Conflict. s View Serializable

1. Same txn must read init. value

2. Same txn must write after reading by same txn's across 2 schedules \rightarrow eg. first read /

3. Same txn for final write | $V.G. = C.G. + \text{Blind then rest writes are writes}$

⑦ \rightarrow ⑦ Precedence graph : conflicting action of T_i occurs before T_j

Recoverable \rightarrow if T_j reads item com. written by $txn T_i$

then commit op. of T_j comes later than T_i .

Cascadeless \rightarrow commit op. of T_i before read of T_j

- 1) Serializable
- 2) Repeatable read
- 3) Read committed
- 4) Uncommitted

Concurrency Control

Two-Phase Locking \rightarrow recoverable

↓ \downarrow serializable ↑

Grow Phase Shrinking Phase

1) Strict \rightarrow hold all x-locks till abort

2) Rigorous \rightarrow hold all locks "

Graph-Based Protocol

Partial ordering $\Rightarrow D = \{d_1, \dots, d_n\}$ only x-locks

\rightarrow database graph

\rightarrow Serializable + No deadlock + shorter wait time

No recoverability or cascade freedom

\rightarrow cannot relock a data

Deadlock prevention: wait (non-pr) \rightarrow older die (c-empty) \rightarrow older wait

Wound: older wounds never

wait younger than waits

Multi-granular: DB \rightarrow area \rightarrow file \rightarrow record

IS: intention-shared / IX / SIX

X \rightarrow false with every other lock

SIX \rightarrow true only with IS info enough

S \rightarrow false with IX to build matrix

Phantom \rightarrow updating / adding reads tuple that satisfy pred

Index-Locking Protocol

\rightarrow lock leaf-nodes

Next-Key Locking Protocol

\rightarrow lock values satisfying lookup

\rightarrow lock next-key value

TSO-Protocol

Read: $TS(T_i) \geq W-t(B)$

$\Rightarrow R_t(B) = \max(R_t(B), TS(T_i))$

else roll-back

Write: $TS(T_i) \geq W-t(B) \wedge \geq R-t(B)$

$\Rightarrow W_t(B) = TS(T_i)$

else roll-back

\Rightarrow serial. + No deadlock

* may not be recoverable \rightarrow som

Timestamp Rule \rightarrow if $TS(T_i) < W-t(B)$

then ignore the write operation

\Rightarrow View-serializable

Validation-Based

1. Read / Ex. phase 2. Valid-phase 3. Write

Validation: $finishTS(T_i) < startTS(T_j)$

$startTS(T_i) < fTS(T_i) < vTS(T_j)$ OR

T_i & written items don't intersect with T_j (read)

MVCC-TSO

$$\max(R_t(B), TS(T_i))$$

Ok: largest t.s. (write) $\leq TS(T_i)$

Read: value of Ok returned $R_t(B) = ?$

Write: $TS(T_i) < R-t(B) \rightarrow$ rolled-back

$TS(T_i) = W-t(B)$: overwritten otherwise \rightarrow new version of Q_i initialised to $TS(T_i)$

MVCC-TPL

$$Wt(Q_i) = 0$$

Update: Follow rigorous TPL \uparrow

First write by T_i results in Q_i [new]

$$TS(T_i) = ts_counter + W-t(Q_i) = TS(T_i)$$

$$ts_c = ts_c + 1$$

no locks needed \rightarrow Read-only:

* $TS(T_i) = ts_counter$ serializable

* Follow MVCC-TPL for reads.

Commit processing into MVCC

Extra + space overhead for drawbacks: tuples

→ Recovery System \rightarrow

logical error (internal) (e.g. deadlocks)

Shadow Copying: $\square \rightarrow \square$ update ptr

$\langle T_i, start \rangle \rightarrow \langle T_i, X, V_i, V_2 \rangle \rightarrow \langle T_i, commit \rangle$

immediate \rightarrow deferred at time modif. → uncommitted modif. of commit

undo $\langle T_i \rangle \rightarrow \langle T_i, X, V \rangle \in \langle T_i, abort \rangle$

Checkpt L list of txns active.

Rollback $\rightarrow \langle T_i, X_i, V_i, V_2 \rangle \Rightarrow \langle T_i, X_j, V_i \rangle$

$\langle T_i, start \rangle \Rightarrow \langle T_i, abort \rangle$

Redo phase: i) checkpt L: set undolist to L

ii) scan forward & redo actions

iii) $\langle T_i, start \rangle \rightarrow$ undo list

iv) $\langle T_i, com/abort \rangle$ - remove from list

Undo phase: Rollback undo-list

No-force policy + steal policy

↳ committed may not be written + uncommitted may be written

* Fuzzy Checkpt → Note list of modified buffer blocks (M).

Output M to disk (don't allow update)

Remote Backup One-Safe \rightarrow Two-stage system \rightarrow Two-very-safe

One-Safe \rightarrow Two-stage system \rightarrow Two-very-safe